

Monte Carlo Reinforcement Learning

Gachon University
Department of AI/SW
202135520 김승수 (Kim SeungSu)




Table of Contents

1. Monte Carlo Method
2. Monte Carlo Method example
3. Background knowledge
4. Monte Carlo Reinforcement Learning
5. ε -greedy exploration pseudo code
6. ε -greedy exploration walkthrough
7. ε -greedy exploration python code and results
8. Monte Carlo Reinforcement Learning Pros and Cons
9. Reference

Monte Carlo Method

“Computational algorithms that rely on repeated random sampling to obtain numerical results”

- Wikipedia

Monte Carlo Method

How to obtain winrate of Solitaire?

1. Calculate odds

possible outcomes = $52! \approx 8 \cdot 10^{67}$

=> not feasible

2. Estimate


play N games and see how many times are wins

=> much easier


Monte Carlo Method example

How to get value of π ?

1. Draw a circle



2. Draw square around circle with length equal to diameter



Monte Carlo Method example

3. Put N points inside the square

- $N = 16$


4. Count points inside circle (k)

- $k = 12$

5. Area of circle / area of square

$$= \pi * r * r / 2r * 2r = \pi / 4$$

$$- 12/16 = 0.75 \approx 0.7853... = \pi / 4$$



Background knowledge

1. Return

$$G_t \doteq \sum_{k=t+1}^T \gamma^{k-t-1} R_k$$

2. Policy

- represented as π
- maps states to probabilities of selecting each possible action
- what we want to learn

$$\max_{\pi} \mathbb{E}_{\pi}[G_t]$$

Background knowledge

3. Value function

- state-value function: expected return when starting in state s and following π

$$v_\pi(s) \doteq \mathbb{E}_\pi[G_t \mid S_t = s] = \mathbb{E}_\pi \left[\sum_{k=0}^{\infty} \gamma^k R_{t+k+1} \mid S_t = s \right], \text{ for all } s \in \mathcal{S}$$

- action-value function: expected return when starting in state s , performing action p and following π

$$q_\pi(s, a) \doteq \mathbb{E}_\pi[G_t \mid S_t = s, A_t = a] = \mathbb{E}_\pi \left[\sum_{k=0}^{\infty} \gamma^k R_{t+k+1} \mid S_t = s, A_t = a \right]$$

Background knowledge

3. Optimal policy

- represented as π^*

4. Optimal state-value function

- represented as v^*

- used to get π^*

$$v_*(s) \doteq \max_{\pi} v_{\pi}(s), \text{ for all } s \in \mathcal{S}$$

5. Optimal action-value function

- represented as q^*

- used to get π^*

$$q_*(s, a) \doteq \max_{\pi} q_{\pi}(s, a), \text{ for all } s \in \mathcal{S} \text{ and } a \in \mathcal{A}(s)$$

Monte Carlo Reinforcement Learning

Characteristics:

1. Learns from episode of experience
2. Model-free
3. Uses idea that value function = empirical mean return
4. When choosing action, randomly choose between exploration and exploitation to prevent local minima

Monte Carlo Reinforcement Learning

Mean of $x_1 \sim x_k$:

$$\begin{aligned}\mu_k &= \frac{1}{k} \sum_{j=1}^k x_j = \frac{1}{k} \left(x_k + \sum_{j=1}^{k-1} x_j \right) \\ &= \frac{1}{k} (x_k + (k-1)\mu_{k-1}) \\ &= \mu_{k-1} + \frac{1}{k} (x_k - \mu_{k-1})\end{aligned}$$

Say $x = g$

$$\mu_k = E(x) = E(g_t \mid S_t = s) = v(S_t)$$

Thus, where $N(S_t)$ represents the amount of time a state was visited:

$$v(S_t) := v(S_t) + 1/N(S_t) * (g_t - v(S_t))$$

This can be used to update the policy

ϵ - greedy exploration pseudo code

For episode in num_episode:

 While not done:

 If rand(0, 1) >= epsilon:

 action = argmax(Q[state, :]) # exploit

 else:

 action = rand(actions) # explore

 new_state,reward,done = action.perform()

 state = new_state

 results_list += (state, action)

 results_sum += reward

 For (state, action) in results_list:

 visit[state, action] += 1

 alpha = 1/visit[state,action]

 Q[state,action] += alpha*(result_sum - Q[state, action])

ε - greedy exploration walkthrough (0)

$\varepsilon = 0.2$

State = x

Actions = u(p), d(own), l(eft),
r(ight)

Starting point	Reward: 0 Done: X	Reward: 0 Done: X
Reward: 0 Done: X	Reward: 10 Done: O	Reward: 0 Done: X
Reward: 0 Done: X	Reward: 0 Done: X	Reward: 100 Done: O

ε - greedy exploration walkthrough (1)

$\varepsilon = 0.2$

State = (0, 0)

Actions = u(p), d(own), l(eft),
r(ight)

rand() = 0.5 $\geq \varepsilon$

argmax(Q[(0,0), :]) = r

Starting point s	Reward: 0 Done: X	Reward: 0 Done: X
Reward: 0 Done: X	Reward: 10 Done: O	Reward: 0 Done: X
Reward: 0 Done: X	Reward: 0 Done: X	Reward: 100 Done: O

ε - greedy exploration walkthrough (2)

$\varepsilon = 0.2$

State = (1, 0)

Actions = u(p), d(own), l(eft),
r(ight)

rand() = 0.1 < ε

rand(actions) = d

Starting point	Reward: 0 Done: X s	Reward: 0 Done: X
Reward: 0 Done: X	Reward: 10 Done: O	Reward: 0 Done: X
Reward: 0 Done: X	Reward: 0 Done: X	Reward: 100 Done: O

ε - greedy exploration walkthrough (3)

$\varepsilon = 0.2$

State = (1, 0)

Actions = u(p), d(own), l(eft),
r(ight)

Done = true

Reward = 10

Starting point	Reward: 0 Done: X	Reward: 0 Done: X
Reward: 0 Done: X	Reward: 10 Done: O s	Reward: 0 Done: X
Reward: 0 Done: X	Reward: 0 Done: X	Reward: 100 Done: O

ϵ - greedy exploration walkthrough (4)

$$\text{Visit}[(0, 0), r] = 0 + 1 = 1$$

$$Q[(0, 0), r] = 0 + 1/1 * (10 - 0) = 10$$

$$\text{visit}[(1, 0), d] = 0 + 1 = 1$$

$$Q[(1, 0), d] = 0 + 1/1 * (10 - 0) = 10$$

Starting point	Reward: 0 Done: X	Reward: 0 Done: X
Reward: 0 Done: X	Reward: 10 Done: O	Reward: 0 Done: X
Reward: 0 Done: X	Reward: 0 Done: X	Reward: 100 Done: O

ε - greedy exploration walkthrough (5)

$\varepsilon = 0.2$

State = (0, 0)

Actions = u(p), d(own), l(eft),
r(ight)

rand() = 0.7 $\geq \varepsilon$

$\text{argmax}(Q[(0,0), :]) = r$

Starting point s	Reward: 0 Done: X	Reward: 0 Done: X
Reward: 0 Done: X	Reward: 10 Done: O	Reward: 0 Done: X
Reward: 0 Done: X	Reward: 0 Done: X	Reward: 100 Done: O

ε - greedy exploration walkthrough (6)

$\varepsilon = 0.2$

State = (1, 0)

Actions = u(p), d(own), l(eft),
r(ight)

rand() = 0.1 < ε

rand(actions) = r

Starting point	Reward: 0 Done: X s	Reward: 0 Done: X
Reward: 0 Done: X	Reward: 10 Done: O	Reward: 0 Done: X
Reward: 0 Done: X	Reward: 0 Done: X	Reward: 100 Done: O

ε - greedy exploration walkthrough (7)

$\varepsilon = 0.2$

State = (2, 0)

Actions = u(p), d(own), l(eft),
r(ight)

rand() = 0.5 $\geq \varepsilon$

argmax(Q[(2,0),:]) = d

Starting point	Reward: 0 Done: X	Reward: 0 Done: X s
	Reward: 0 Done: X	Reward: 10 Done: O
	Reward: 0 Done: X	Reward: 100 Done: O

ε - greedy exploration walkthrough (8)

$\varepsilon = 0.2$

State = (2, 1)

Actions = u(p), d(own), l(eft),
r(ight)

rand() = 0.9 $\geq \varepsilon$

argmax(Q[(2,1),:]) = d

Starting point	Reward: 0 Done: X	Reward: 0 Done: X
Reward: 0 Done: X	Reward: 10 Done: O	Reward: 0 Done: X S
Reward: 0 Done: X	Reward: 0 Done: X	Reward: 100 Done: O

ε - greedy exploration walkthrough (9)

$\varepsilon = 0.2$

State = (2, 2)

Actions = u(p), d(own), l(eft),
r(ight)

Done = 100

Reward = 100

Starting point	Reward: 0 Done: X	Reward: 0 Done: X
	Reward: 0 Done: X	Reward: 10 Done: O
	Reward: 0 Done: X	Reward: 100 Done: O S

ϵ - greedy exploration walkthrough (10)

$$\text{visit}[(0, 0), r] = 1+1 = 2$$

$$Q[(0, 0), r] = 10 + 1/2*(100-10) = 55$$

$$\text{visit}[(1, 0), r] = 0 + 1 = 1$$

$$Q[(1, 0), r] = 0 + 1/1*(100-0) = 100$$

$$\text{visit}[(2, 0), d] = 0 + 1 = 1$$

$$Q[(2, 0), d] = 0 + 1/1*(100-0) = 100$$

$$\text{visit}[(2, 1), d] = 0 + 1 = 1$$

$$Q[(2, 1), d] = 0 + 1/1*(100-0) = 100$$

Starting point	Reward: 0 Done: X	Reward: 0 Done: X
Reward: 0 Done: X	Reward: 10 Done: O	Reward: 0 Done: X
Reward: 0 Done: X	Reward: 0 Done: X	Reward: 100 Done: O

ϵ - greedy exploration python code (1)

```
import numpy as np
import random
import math
import matplotlib.pyplot as plt

def random_action(actions, state, row_size, column_size):
    while(True):
        # choose random action
        action = random.choice(actions)

        # make sure action is possible
        if action == "up":
            if state[0] > 0:
                return 0
        elif action == "down":
            if state[0] < row_size - 1:
                return 1
        elif action == "left":
            if state[1] > 0:
                return 2
        elif action == "right":
            if state[1] < column_size - 1:
                return 3

def perform_action(action, state, rewards):
    new_state = state.copy()
    if action == 0:
        new_state[0]-=1
    elif action == 1:
        new_state[0]+=1
    elif action == 2:
        new_state[1]-=1
    elif action == 3:
        new_state[1]+=1

    reward = got_reward(new_state, rewards)
    if reward != 0:
        return new_state, reward, True
    else:
        return new_state, 0, False

def got_reward(state, rewards):
    for reward in rewards:
        if (state[0],state[1]) == reward[0]:
            return reward[1]

    return 0

def visualize_path(results_list, row_size, column_size):
    grid = np.zeros((row_size, column_size))
    for state, action in results_list:
        grid[state[0]][state[1]] = action + 1

    print("=-100")
    print(len(results_list))
    for r in range(row_size):
        for c in range(column_size):
            if grid[r][c] == 0:
                print("  ", end = "")
            elif grid[r][c] == 1:
                print("v ", end = "")
            elif grid[r][c] == 2:
                print("^ ", end = "")
            elif grid[r][c] == 3:
                print("< ", end = "")
            elif grid[r][c] == 4:
                print("> ", end = "")
            print("")

    print("=-100")
    actions = ["up", "down", "right", "left"]
    # parameters =====
    row_size = 5
    column_size = 5
    epsilon = 0.3
    num_episodes = 20000
    rewards = [(1, 3), ((0,3), 70), ((4,4),100000), ((2, 3), 5000)]
    # =====
    avg_reward_per_action = [0]*5
    episode_count = [0]*5

    Q = np.zeros((row_size,column_size,len(actions)))
    visit_count = np.zeros((row_size, column_size, len(actions)))

    for episode in range(num_episodes):
        # reset variables
        state = [0, 0]
        result = 0
        results_list = []
        results_sum = 0
        reward = 0
        done = False
        count = 0

        # complete episode
        while not done:
            if np.random.rand() > epsilon:
                action = np.argmax(Q[state[0], state[1], :])
                # if max is 0, choose randomly
                if Q[state[0], state[1], action] == 0:
                    action = random_action(actions, state, row_size, column_size)
            else:
                action = random_action(actions, state, row_size, column_size)

            new_state, reward, done = perform_action(action, state, rewards)
            results_list.append((state, action))
            state = new_state.copy()
            results_sum += reward
            count+=1

        # save values
        index = math.floor(math.log10(episode+1))
        episode_count[index]+=1
        avg_reward_per_action[index] += 1/episode_count[index]*(reward/count - avg_reward_per_action[index])
        #print(f'{episode}: {reward/count}')

        # update Q matrix
        for (state, action) in results_list:
            visit_count[state[0], state[1], action] += 1
            alpha = 1/visit_count[state[0], state[1], action]
            Q[state[0], state[1], action] += alpha*(results_sum - Q[state[0], state[1], action])

        # print path
        visualize_path(results_list, row_size, column_size)

    # print results
    for i in range(len(avg_reward_per_action)):
        print(f'range: -10^{i+1} - average reward per action: {avg_reward_per_action[i]:.3f} data size: {episode_count[i]}')
        #print(f'{avg_reward_per_action[i]}, {episode_count[i]}')

    # visualize results
    plt.bar(["1-9", "10-99", "100-999", "10000-99999", "100000-999999"], avg_reward_per_action)
    plt.show()
```

ϵ - greedy exploration python code (2)

Current version (my Google drive):

[https://drive.google.com/file/d/1GcvEHfC8JNfFg516VT7YL4QQbFkZBfFw/view
?usp=sharing](https://drive.google.com/file/d/1GcvEHfC8JNfFg516VT7YL4QQbFkZBfFw/view?usp=sharing)

Newest version (my Github):

<https://github.com/DeceptiveRat/Monte-Carlo-Reinforcement-Learning>

ϵ - greedy exploration python results (1)

At first, the agent goes to nearby rewards and takes long paths




Later, the agent takes short paths to rewards far away



ϵ - greedy exploration python results (2)

```
range: ~10^1 - average reward per action: 160.338 data size: 9
range: ~10^2 - average reward per action: 497.321 data size: 90
range: ~10^3 - average reward per action: 4861.334 data size: 900
range: ~10^4 - average reward per action: 6736.648 data size: 9000
range: ~10^5 - average reward per action: 7657.695 data size: 10001
```

Average reward per action can be seen going up as episodes pass



Pros and Cons

Pros:

- doesn't require a model
- conceptually simple
- easy to implement

Cons:

- inefficient for long episodes
- requires episodes to be finished to learn from it

References

1. Durham University Lectures:

<https://www.google.com/url?sa=t&source=web&rct=j&opi=89978449&url=https://cwkx.github.io/data/teaching/dl-and-rl/rl-lecture5.pdf&ved=2ahUKEwi99cWsoaiQAxWCsVYBHQtuNhwQFnoECB4QAQ&usg=AOvVaw1OS2PUTCnzYGNTCOOr3Ecks>

2. Wikipedia:

https://en.wikipedia.org/wiki/Monte_Carlo_method

3. Medium:

<https://medium.com/@hsinhungw/intro-to-reinforcement-learning-monte-carlo-to-policy-gradient-1c7ede4eed6e>

**THANK
YOU**